

Introducing the Concept of Expectation



Dmitri Nesteruk

QUANTITATIVE ANALYST

@dnesteruk <http://activemesa.com>



Goals:

Understand the notion of an
expectation and associated
statistical measures.



Overview



Expectation

Mean

Law of the Unconscious Statistician

Variance

Moments and the Moment Generating Function

Joint Distributions

Covariance and Correlation



Expectation

If you measure a newborn baby, how tall do you expect it to be?

We don't know, since it's a random variable

Measure n newborn babies (the more the better) and average out the result

$$\bar{x} = \frac{\sum_{i=1}^n x_i}{n}$$

This is the *mean* or *average* baby height

So, given a random variable X representing a baby's height, \bar{x} is its *expected value* (51cm/20in, actually)



Expectation of a Random Variable

The expectation (expected value, mean, average) of a random variable is a weighted average of all possible values weighted by their respective probabilities.



Expectation Example

Suppose a stock price can increase by 1 point with $P_U = 0.6$, stay in place with $P_S = 0.1$ and go down by 1 point with $P_D = 0.3$

If X represents the change in stock price, what is $E[X]$?

$$E[X] = 1 \cdot 0.6 + 0 \cdot 0.1 + 0.3 \cdot (-1) = 0.3$$

In other words, on average, we'd expect the stock price to go up by 0.3 points



Mean

For a bounded discrete random variable X with probability function f , the expectation of X is

$$E[X] = \sum_x x f(x)$$

Also referred to as the *mean* of X or the *expected value* of X .



Mean of a Bernoulli Random Variable

A rv X has a Bernoulli distribution when X can only take values 0 or 1 with $P(X = 1) = p$

$$E[X] = 0 \cdot (1 - p) + 1 \cdot p = p$$

Note that the expectation does not have to be a possible value of the variable



Existence of Mean

Given a variable X with a pf f

Consider

$\sum_{\text{Positive } x} xf(x)$ and $\sum_{\text{Negative } x} xf(x)$

If at least one of these is finite, then $E[X]$ is said to exist, defined as

$$\sum_{\text{All } x} xf(x)$$

If both of the sums are infinite, mean does not exist



Existence of Mean Example

Let X be a rv with pf

$$f(x) = \begin{cases} \frac{1}{2|x|(|x|+1)}, & x = \pm 1, \pm 2, \pm 3, \dots \\ 0, & \text{otherwise} \end{cases}$$

Check the sums

$$\sum_{-1}^{-\infty} x \frac{1}{2|x|(|x|+1)} = -\infty$$

$$\sum_1^{\infty} x \frac{1}{2|x|(|x|+1)} = \infty$$

Both sums are infinite $\therefore E[X]$ does not exist



Expectation for a Continuous Distribution

For a continuous bounded random variable,
a sum turns into an integral

If X is a bounded continuous rv with pdf f ,

$$E[X] = \int_{-\infty}^{\infty} xf(x) dx$$

also called the *expected value* or *mean*



Continuous Expectation Example

Suppose the time X until a part fails has the pdf

$$f(x) = \begin{cases} 2x, & 0 < x < 1 \\ 0, & \text{otherwise} \end{cases}$$

Then

$$E[X] = \int_0^1 x \cdot 2x \, dx = \int_0^1 2x^2 \, dx = \frac{2}{3}$$

The expectation of the distribution with pdf f is $2/3$



Existence of Continuous Mean

If either $\int_0^{\infty} xf(x) dx$ or $\int_{-\infty}^0 xf(x) dx$ is finite,

The mean/expectation/expected value of X exists and is defined as

$$\int_{-\infty}^{\infty} xf(x) dx$$

If both integrals are infinite, then the mean does not exist



Existence of Continuous Mean Example

Consider the Cauchy distribution with pdf

$$f(x) = \frac{1}{\pi(1+x^2)}$$

This is a pdf, i.e. $\int_{-\infty}^{\infty} f(x) = 1$ because

$$\frac{d(\tan^{-1} x)}{dx} = \frac{1}{1+x^2}$$

Calculate the two integrals:

$$\int_0^{\infty} \frac{x}{\pi(1+x^2)} dx = \infty \text{ and } \int_{-\infty}^0 \frac{x}{\pi(1+x^2)} dx = -\infty$$

Thus the mean does not exist

Function of a Random Variable

If X is a rv with pdf f , consider a real-valued function $r(X)$

To find the expected value of $r(X)$, we

- Let $Y = r(X)$
- Determine probability distribution of Y
- Determine $E[Y]$

$E[r(x)] = E[Y] = \int_{-\infty}^{\infty} yg(y) dy$
if the expectation exists

Expectation of Function Example

Given a rv X with pdf of $f(x) = 3x^2, 0 < x < 1$

Let $r(x) = 1/x$

First we calculate pdf of $Y = r(X)$:

$$g(y) = 3y^{-4}, \quad y > 1$$

This gives us the mean

$$E[Y] = \int_0^{\infty} y \cdot 3y^{-4} dy = \frac{3}{2}$$

These manipulations are unnecessary due to LOTUS



Law of the Unconscious Statistician

If X is a rv and $r(x)$ is a real-valued function of a real variable then, if X has a discrete distribution,

$$E[r(x)] = \sum_{All\ x} r(x)f(x)$$

If X has a continuous distribution,

$$E[r(x)] = \int_{-\infty}^{\infty} r(x)f(x) dx$$

Expectation of Function Example (LOTUS)

Given a rv X with pdf of $f(x) = 3x^2, 0 < x < 1$

Let $r(x) = 1/x$

$$E[Y] = \int_0^1 r(x)f(x) dx = \int_0^1 \frac{1}{x} \cdot 3x^2 dx = 3/2$$

Note: in general, $E[f(X)] \neq f(E[X])$



Properties of Distributions

If X is a rv for which $E[X]$ exists,

If $X = c$ with probability 1, $E[X] = c$

If $Y = aX + b$ where a and b are finite constants (linear function),

$$E[Y] = aE[X] + b$$

If X_1, \dots, X_n are n rvs s.t. $E[X_i]$ is finite, then

$$E[X_1 + \dots + X_n] = E[X_1] + \dots + E[X_n]$$

$$E[X_1 \times \dots \times X_n] = E[X_1] \times \dots \times E[X_n]$$



Variance

A mean does not completely describe a distribution

A uniform distribution from -1 to 1 and a standard normal distribution have same mean, but are drastically different

We need another measure of how 'spread out' the distribution is



Variance

If X is a rv with finite mean $\mu = E[X]$, the *variance* of X is defined as

$$V[X] = E[(X - \mu)^2]$$

If X has infinite mean or if $E[X]$ does not exist, $V[X]$ does not exist.

The *standard deviation* of X is the non-negative square root of $V[X]$ if the variance exists:

σ – standard deviation, σ^2 – variance



Variance Example

Consider a coin flip that yields a 0 or a 1 with equal probability $P = 0.5$

$$\mu = E[X] = 0 \cdot 0.5 + 1 \cdot 0.5 = 0.5$$

Define $Y = (X - \mu)^2$. Then $V[X] = E[Y]$

Compute pf of Y :

x	0	1
y	0.25	0.25
$f(y)$	0.5	0.5

$$V[X] = E[Y] = 0.5(0.25 + 0.25) = 0.25$$

$$\text{Standard deviation } \sigma = \sqrt{0.25} = 0.5$$



Alternative Variance Calculation Method

We can expand variance calculation

$$\begin{aligned} V[X] &= E[(X - \mu)^2] \\ &= E[X^2 - 2\mu X + \mu^2] \\ &= E[X^2] - 2\mu E[X] + \mu^2 \\ &= E[X^2] - \mu^2 \end{aligned}$$

In other words,

$$V[X] = E[X^2] - (E[X])^2$$

Additive Property of Variance

Given independent random variables X_1, \dots, X_n with finite means and constants a_1, \dots, a_n and b ,
 $V[a_1X_1 + \dots + a_nX_n + b] = a_1^2V[X_1] + \dots + a_n^2V[X_n]$

Notice how b does not contribute to the variance



Moments

For a random variable X and $k \in \mathbb{N}$, the expectation $E[X^k]$ is called the k th moment of X

The mean $E[X]$ is the first moment of X

The k th moment exists iff $E[|X|^k] < \infty$



Central Moment

An n th central moment μ_n is the moment of a probability distribution about the mean

$$\mu_n = E[(X - E[X])^n] = \int_{-\infty}^{\infty} (x - \mu)^n f(x) dx$$

$$\mu_0 = 1 \text{ (because } \int 1 \cdot f(x) dx = 1)$$

$$\mu_1 = 0 \text{ (because } E[X - E[X]] = E[X] - E[E[X]] = 0)$$

μ_2 is the variance



Moment Generating Function

Given a rv X , for each $t \in \mathbb{R}$, the *moment generating function* (mgf) of X is defined as

$$\psi(t) = E[e^{tX}]$$

For $n \in \mathbb{N}$, the n th moment of X is the n th derivative of the mgf at $t = 0$, i.e.,

$$E[X^n] = \psi^{(n)}(0)$$

Calculating an MGF

Suppose X has a pdf $f(x) = e^{-x}$ for $x > 0$

Let's find the mgf:

$$\begin{aligned}\psi(t) &= E[e^{tX}] = \int_0^{\infty} e^{tx} e^{-x} dx \\ &= \int_0^{\infty} e^{(t-1)x} dx\end{aligned}$$

The integral diverges for $t \geq 1$, so mgf is

$$\psi(t) = \frac{1}{1-t}$$

for $t < 1$



Calculating variance from MGF

Mgf is $\psi(t) = \frac{1}{1-t}$

Calculate derivatives:

$$\psi'(t) = \frac{1}{(1-t)^2} \text{ and } \psi''(t) = \frac{2}{(1-t)^3}$$

Armed with this, we can calculate variance:

$$\begin{aligned} E[X^2] - (E[X])^2 &= \psi''(0) - [\psi'(0)]^2 \\ &= 2 - 1 = 1 \end{aligned}$$

Thus, $V[X] = 1$



Means and Variances of Some Distributions

Distribution	Pmf	Mean	Variance
Uniform (disc)	$\frac{1}{b-a+1}$ for $x = a, \dots, b$	$\frac{a+b}{2}$	$\frac{(b-a+1)^2 - 1}{12}$
Binomial	$\binom{n}{k} p^k (1-p)^{n-k}$	np	$np(1-p)$
Geometric	$(1-p)^k p$	$\frac{1-p}{p}$	$\frac{1-p}{p^2}$
Hypergeometric	$p(x) = \frac{\binom{r}{x} \binom{N-r}{n-x}}{\binom{N}{n}}$	$n \frac{r}{N}$	$n \frac{r}{N} \frac{(N-r)}{N} \frac{N-n}{N-1}$

Means and Variances of Some Distributions

Distribution	Pdf	Mean	Variance
Uniform (cont)	$\frac{1}{b-a}$	$\frac{a+b}{2}$	$\frac{(b-a)^2}{12}$
Normal	$\frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$	μ	σ^2
Gamma	$\frac{x^{\alpha-1} e^{-\frac{x}{\beta}}}{\beta^\alpha \Gamma(\alpha)}$	$\frac{\alpha}{\beta}$	$\frac{\alpha}{\beta^2}$
Beta	$\frac{x^{\alpha-1} (1-x)^{\beta-1}}{B(\alpha, \beta)}$	$\frac{\alpha}{\alpha+\beta}$	$\frac{\alpha\beta}{(\alpha+\beta)^2(\alpha+\beta+1)}$



Joint Distributions

Interested in several random variables related to each other

E.g., if you pick a random family and measure their earnings and how many cars they have, those things *might* be related

Instead of one variable, we consider two or more

We'll focus on 2 rvs



Joint Distribution Example

Suppose we roll *two* dice

We want the probability that first dice is a 1
and second is a 3

$$P(X = 1, Y = 3) = \frac{1}{36}$$

A general calculation can be expressed as a
joint pmf



Joint Probability Mass Function

Given two discrete random variables X and Y , their joint probability mass function (pmf) is defined as

$$P_{XY}(x, y) = P(X = x, Y = y)$$

The joint pmf need to satisfy

$$\sum_{x \in X, y \in Y} P_{XY}(x, y) = 1$$

Marginal PMF

The joint pmf has all the information regarding the distribution of X and Y

We can try to ‘extract’ the pmf of X from P_{XY} :

$$P_X(x) = P(X = x) = \sum_{y_i} P(X = x, Y = y_i)$$

We call $P_X(x)$ the *marginal* pmf of X .



Marginal PMF Example

Consider rvs X and Y with joint pmf

	$Y = 0$	$Y = 1$	$Y = 2$
$X = 0$	$1/6$	$1/4$	$1/8$
$X = 1$	$1/8$	$1/6$	$1/6$

Let's find $P_X(x)$

- $P_X(0) = \frac{1}{6} + \frac{1}{4} + \frac{1}{8} = 13/24$
- $P_X(1) = 1 - P_X(0) = 11/24$

$$P_X(x) = \begin{cases} \frac{13}{24}, & x = 0 \\ \frac{11}{24}, & x = 1 \\ 0, & \text{otherwise} \end{cases}$$

Joint Cumulative Distribution Function

The joint cumulative distribution function of two random variables X and Y is defined as

$$F_{XY}(x, y) = P(X \leq x, Y \leq y)$$

We expect that

$$\begin{aligned} F_{XY}(\infty, \infty) &= 1 \\ F_{XY}(-\infty, y) &= 0, & \text{for any } y \\ F_{XY}(x, -\infty) &= 0, & \text{for any } x \end{aligned}$$



Functions of 2 or More Random Variables

Sometimes we have a function which involves two or more rvs

- E.g., sum of two dice being rolled
- $r(X, Y)$

LOTUS for two discrete rvs

$$E[r(X, Y)] = \sum_{x_i, y_i} r(x_i, y_i) P_{XY}(x_i, y_i)$$

Joint Probability Density Function

Two random variables X and Y are *jointly continuous* if there exists a non-negative function f_{XY} s.t. for any set $A \in \mathbb{R}^2$,

$$P((X, Y) \in A) = \iint_A f_{XY}(x, y) \, dx \, dy$$

The function f_{XY} is called the *joint probability density function* of X and Y . Naturally,

$$\int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f_{XY}(x, y) \, dx \, dy = 1$$

Joint PDF Example

Suppose X and Y have joint pdf

$$f_{XY}(x, y) = x + cy^2, \quad 0 \leq x \leq 1, 0 \leq y \leq 1$$

Let's find the constant c

$$\begin{aligned} \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f_{XY}(x, y) dx dy &= \int_0^1 \int_0^1 x + cy^2 dx dy \\ &= \left[\frac{y}{2} + \frac{cy^3}{3} \right]_{y=0}^{y=1} = \frac{1}{2} + \frac{1}{3}c = 1 \end{aligned}$$

Therefore, $c = 3/2$

Marginal PDFs

Similarly to marginal pmfs, marginal pdfs involve leaving one of the two variables alone, i.e.

$$f_X(x) = \int_{-\infty}^{\infty} f_{XY}(x, y) dy, \quad \text{for all } x$$
$$f_Y(y) = \int_{-\infty}^{\infty} f_{XY}(x, y) dx, \quad \text{for all } y$$

This lets us find a pdf for one of the two (or more) variables from a joint pdf



Marginal PDF Example

Given the joint pdf

$$f(x, y) = 10x^2y, \quad 0 \leq y \leq x \leq 1$$

Let's find $f_X(x)$ and $f_Y(y)$

- Note the domains of x and y !

$$f_X(x) = \int_{-\infty}^{\infty} f_{XY}(x, y) dy = \int_0^x 10x^2y dy = 5x^4$$

$$f_Y(y) = \int_y^1 10x^2y dx = \frac{10}{3}y(1 - y^3)$$

Search for Related Variables

Does smoking cause cancer?

Need to take measurements of cancer and amounts smoked

- Also variables that influence the result

Do the two sets of values vary together (as one grows, so does another)?

- Calculate covariance

Can we get a standardized measure of how related they are?

- Calculate correlation

Reminder: correlation does not imply causation



Covariance

If X and Y are rvs with finite means $E[X] = \mu_X$ and $E[Y] = \mu_Y$, the covariance of X and Y is defined as

$$\text{cov}(X, Y) = \sigma_{XY} = E[(X - \mu_X)(Y - \mu_Y)]$$

if the expectation exists.

Covariance Example

Let X, Y have a joint pdf

$$f(x) = \begin{cases} 2xy + 0.5, & 0 \leq x \leq 1, 0 \leq y \leq 1 \\ 0, & \text{otherwise} \end{cases}$$

Let's compute $\text{cov}(X, Y)$

First we need $\mu_X = \mu_Y$

$$\mu_X = \int_0^1 \int_0^1 (2x^2y + 0.5) dy dx = 7/12$$

Then,

$$\begin{aligned} & \text{cov}(X, Y) \\ &= \int_0^1 \int_0^1 \left(x - \frac{7}{12}\right) \left(y - \frac{7}{12}\right) (2xy + 0.5) dy dx \\ &= 1/144 \end{aligned}$$

What is Covariance?

Covariance gives a numerical measure to the degree to which X and Y vary together

The magnitude of $\text{cov}(X, Y)$ is influenced by X and Y

Correlation is a much better, standardized measure



Correlation

Given random variables X and Y with finite variances σ_X^2 and σ_Y^2 , the *correlation* of X and Y is defined as

$$\rho(X, Y) = \frac{\text{cov}(X, Y)}{\sigma_X \sigma_Y}$$



Correlation Values

Correlation values range from -1 to 1

$\rho = 0$ means variables are uncorrelated (we expect this for independent variables)

$\rho = 1$ means as one variable grows/shrinks, so does another, perfectly

- Correlation measures only a linear relationship

$\rho = -1$ means variables have a perfect inverse relationship (as one grows, another shrinks)



Summary



A random variable has an expected value $E[X]$ also called mean or average

Variance ($E[(X - \bar{x})^2]$) measures how 'spread out' the values are

Two or more variables can have a joint pmf/pdf/cdf

- Marginal functions can be extracted from joint function

Covariance/correlation are measures of a linear relationship between two variables